Enterprise Artificial Intelligence Governance & Algorithmic Ethics Policy

Enterprise Compliance Directive • Ref: POL-GOV-004

ORGANIZATIONAL POLICY MANUAL	**Policy Code:** POL-GOV-004
Document: Enterprise Artificial Intelligence Governance & Algorithmic Ethics Policy	**Version:** 2.0 | **Effective:** 2025-01-20
Category: AI Governance & Ethics	**Business Area:** Product, Data Science & AI Engineering
Compliance Owner: Chief AI Ethics Officer & Model Risk Committee	**Classification:** Confidential - Internal Corporate Use

1. Policy Purpose & Executive Scope

Establishes compliance frameworks for AI systems, generative AI adoption, algorithmic bias prevention, human-in-the-loop controls for high-risk applications, and data lineage tracking.

2. Binding Regulatory Clauses & Compliance Requirements

Section 1.1: High-Risk AI System Classification & Risk Register [POL-GOV-004-C01]

Requirement: MANDATORY | Risk Impact: CRITICAL | Regulatory Base: EU Artificial Intelligence Act (Regulation 2024/1689) Annex III; NIST AI RMF 1.0

Any artificial intelligence, machine learning, or automated scoring system used to make or substantially influence decisions concerning creditworthiness, employment screening, customer identity verification, healthcare recommendations, or insurance pricing is classified as a High-Risk AI System. Such systems must be registered in the Corporate AI Registry and undergo formal Model Risk Committee (MRC) approval prior to deployment.

Section 2.3: Human-in-the-Loop (HITL) for Automated Customer Decisions [POL-GOV-004-C02]

Requirement: MANDATORY | Risk Impact: HIGH | Regulatory Base: GDPR Article 22 (Automated Decision Making); EU AI Act Article 14

Fully autonomous decisions that produce legal effects or significantly affect individuals (such as loan rejections, account terminations, or hiring disqualifications) are strictly prohibited without meaningful Human-in-the-Loop (HITL) oversight. A qualified human reviewer must possess the operational authority, capability, and time to understand the algorithmic rationale and override the recommendation.

Section 3.2: Algorithmic Bias Testing & Disparate Impact Disparity Ratio [POL-GOV-004-C03]

Requirement: MANDATORY | Risk Impact: HIGH | Regulatory Base: EEOC Uniform Guidelines on Employee Selection Procedures; FTC Fair Lending AI Guidance

Prior to release and quarterly thereafter, all predictive models affecting end-users must undergo demographic parity and disparate impact testing across protected characteristics (race, gender, age, religion). The disparate impact ratio between any protected group and the baseline group must not fall below 0.80 (four-fifths rule). Any model breaching this threshold

must be suspended immediately.

Section 4.1: Generative AI & LLM Input Data Confidentiality [POL-GOV-004-C04]

Requirement: MANDATORY | Risk Impact: HIGH | Regulatory Base: Enterprise Trade Secret Protection; GDPR Art. 5; ISO 42001

Employees and engineering systems are prohibited from entering non-public corporate intellectual property, source code, customer personal data (PII), or confidential contract terms into external public generative AI platforms (e.g., public consumer web chatbots) unless an enterprise agreement with strict zero-data-retention and zero-model-training clauses is officially executed.

Section 5.2: Model Explainability & Adverse Action Notices [POL-GOV-004-C05]

Requirement: MANDATORY | Risk Impact: MEDIUM | Regulatory Base: Equal Credit Opportunity Act (ECOA) 12 CFR 1002.9; FCRA § 615(a)

When an AI model produces an adverse outcome against a consumer, the system must generate a human-readable explanation specifying the top 4 contributing feature factors (e.g., using SHAP or LIME attributions). The affected individual must be provided with an Adverse Action Notice within 30 days clearly stating the key reasons for the outcome.